

CS 475 Machine Learning: Homework 3

Deep Learning

Due: Thursday November 7, 11:59pm

100 Points Total Version 1.0

Make sure to read from start to finish before beginning the assignment.

1 Programming (100 points)

There are no analytical question in this assignment.

1.1 Overview

In this assignment you will learn to use PyTorch (<http://pytorch.org/>) to train classifiers to recognize simple images of clothing articles using a subset of the FashionMNIST dataset (<https://github.com/zalandoresearch/fashion-mnist>).

Your grade will be based on a mixture of a) performance on held-out test data of 3 models which you will implement, and b) completing the visualization exercises in the Jupyter notebook.

Requirements You will need Python 3.6. The Python libraries that you will need (and are allowed to use) are given to you in the `requirements.txt` file. Install them using `pip3: pip3 install -r requirements.txt`.

Suggestions Since some of these experiments will take some time to run, and will eat up CPU usage, you might prefer to run things on the Computer Science undergraduate or graduate grid, including the Jupyter notebook. However, this is not required as you should be able to run everything on your own computer. Running on the grid will make it easier to run multiple experiments in parallel.

Additionally, since we will be exploring many parameter settings for each model, you may want to store the models while you explore parameters, and then turn in the model with the best performance. This way you don't need to retrain a model once you've decided on the best hyperparameters for that model.

1.2 Introduction to PyTorch

PyTorch documentation can be found at <http://pytorch.org/docs/stable/index.html>. There are numerous PyTorch resources online.

While it is not required, we strongly recommend that you go through this tutorial: http://pytorch.org/tutorials/beginner/deep_learning_60min_blitz.html.

- You may not have access to a GPU or a CUDA-enabled environment. This is okay; just modify these parts of the tutorial to run on the CPU, which should be fine since the tutorial is not computationally heavy.

- The point of this tutorial is to familiarize you with PyTorch, and in particular to familiarize you with the `torch.nn` and `torch.optim` packages. We suggest that you go through the examples from scratch, modifying parts whenever you are confused.

1.3 Data

The data we will use for this assignment is a subset of the FashionMNIST dataset. Our version consists of 70,000 images (7,000 test, 7,000 dev, 56,000 train). Each image is a 28x28 grayscale article of clothing, and is labeled out of 10 classes. The classes are Shirt, Sneaker, Bag, Ankle Boot, T-shirt, Trouser, Pullover, Dress, Coat, Sandal (with label number in that order).

Read more here: <https://github.com/zalandoresearch/fashion-mnist>. Note, though, that the order of our labels doesn't match the order on the website.

The dataset is provided as 5 separate `.numpy` files. In particular, you should have `train.feats.npy`, `train.labels.npy`, `dev.feats.npy`, `dev.labels.npy`, and `test.feats.npy`.

You will train your models on the `train` data, and evaluate it on the `dev` data. As in past assignments, the unlabeled `test` data is just there for you to make sure you can run your model on it without crashing. Ultimately, we will be running your models on `test` data (with actual labels), and evaluating your model's output.

The data loader we have provided expects all these files to be in the same directory, and loads them by name, so if you rename these files or move them into separate directories, be sure to change the data loader, or update the way you load data.

1.4 Command Line Arguments

The code we provide is a general framework for running pytorch models. It has the following generic arguments:

- **mode**: A non-optional argument. It is either `train` or `predict`. This must be the first argument you pass when running `main.py`.
- **--data-dir**: Should point to the directory where your data files are stored (see the above section).
- **--log-file**: Points to the location where a `csv` file containing logs from your model's training should be stored (only used during `train` mode).
- **--model-save**: Points to the location where your model will be stored after training (during `train` mode) or where the stored model should be loaded from (during `test` mode).
- **--predictions-file**: Points to where to output model predictions (only used during `test` mode).

Do not change these arguments.

Specifically, when we run your models on test data, the command will be (for instance on a saved `ff` model):

```
python3 main.py predict --data-dir ./data/ --model-save ff.torch
--predictions-file ff-preds
```

Before you submit a saved model, make sure you can run this command (with the correct data directory) and that it outputs valid test predictions.

`main.py` also has the following generic hyperparameters (hyperparameters used in the training of all models):

- `--model`: This is the type of model to train. There are 3 models we will build: `simple-ff`, `simple-cnn`, and `best`.
- `--train-steps`: The number of steps to train on. Each step trains on one batch.
- `--batch-size`: The number of examples to include in your batch.
- `--learning-rate`: The learning rate to use with your optimizer during training.

While you should not modify these parameters, you may add to them to experiment with different models. You will then need to modify the `train` function in `main.py`. You can change the training loop to make it better. As long as we can load your stored model using your `test` method, you can train it however you'd like.

The file has a few *model-specific* arguments. You will likely need to add to these when building your `best` model. The feed forward model has *one* model-specific hyperparameter:

- `--ff-hunits`: The number of hidden units in feed-forward layer 1.

And the CNN model has *three* model-specific hyperparameters:

- `--cnn-n1-channels`: The number of channels in the first conv layer.
- `--cnn-n1-kernel`: The size of the square kernel of the first conv layer.
- `--cnn-n2-kernel`: The size of the square kernel of the second conv layer.

What values for these parameters, and what additional parameters you add are up to you!

If you add additional parameters, they should be **optional parameters only**. If we run your train method, we will only provide the parameters listed above. Any other parameters you add must be optional to avoid an error.

1.5 Jupyter Notebook Visualization (35 points)

We will use a Jupyter notebook to visualize different aspects of NN training.

You should start this homework in the Jupyter Notebook we provide. You will be asked to walk through some initial data visualization cells, so that you can become familiar with the data. Eventually you will need to return to your code, and implement a model in order to continue using the notebook.

The notebook serves two purposes. First, the notebook will help you debug your model while you are building it. That is, each model has a portion of the notebook devoted to plotting loss and accuracy, as well as doing hyperparameter sweeps. Once you have implemented those cells, you can use them to help you further debug your models and make decisions about hyperparameter values.

Make sure that you have worked through the entire notebook and run all cells before you finish. We will grade your notebook.




Figure 1: A diagram of the Linear model architecture.

1.6 Implementing a Simple Two-Layer NN (15 points)

You will implement a simple model to classify images of size $28 \times 28 = 784$ pixels from the FashionMNIST dataset into one of the 10 classes. You will use the `Sequential` module to map each input image to hidden layers and eventually to a vector of length 10, which contains *class scores* or *logits*. Since the data comes in the form of vectors of length 784, you can pass it in as is. You may recall logits from homework 1, where we were also doing multiclass classification. **The output of your models will be these scores.**

These scores will be passed directly to a loss function. In this homework, we will use cross-entropy as our loss function. This is exactly the same loss as we used for our multiclass logistic regression model in homework 1. **You will *implicitly* rely on the softmax function through PyTorch's cross-entropy loss function however, so you never need to use the softmax function directly.**

1. Complete the `SimpleFF` class in `models.py` by implementing the `forward` function. (See Figure 1) This model uses one linear layer to map from the inputs to n hidden units, with ReLU activations, and uses another linear layer to map from the hidden units to vectors of length 10. You can apply relu activations to the output of the first layer with `torch.nn.functional.relu`.
2. When you have finished implementing the model, train it via `main.py`. An example command might be: `python code/main.py train --data-dir data --log-file logs/ff-logs.csv --model-save models/ff.torch --model simple-ff`
3. Complete sections 2.1 and 3.1 of the notebook. You might use these blocks to help you debug the model, if you are not getting good development accuracy (around 80%).
4. Save a your version of the model which uses default hyperparameters for all values except for learning rate. Use the best learning rate you found from section 3.1 as the learning rate for the model you submit. Name the model you submit `ff.torch`. Run this model on the test data using `predict` mode, and name the predictions `ff-predictions.txt`. **Submit both the model file and the predictions file.**




Figure 2: A simple diagram of the CNN model architecture.

1.7 Implementing a Simple Convolutional NN (20 points)

The next step is to create a convolutional neural network for multiclass classification, again using cross-entropy loss. As a reminder, 2-D convolutional neural networks operate on *images* rather than *vectors*. (This means that your model will need to reshape the input into a 2 dimensional image!) Their key property is that hidden units are formed using local spatial regions of the input image, *with weights that are shared over spatial regions*. To see an animation of this behavior, see https://github.com/vdumoulin/conv_arithmetic. In particular, pay attention to the first animation (no padding, no strides) and to the 9th animation (no padding, with strides). Here, a 3×3 convolutional filter is run over the image, with the image shown in blue and the output map (which is just a collection of the hidden units) shown in green. You can think of each convolutional filter as a small, translation-invariant pattern detector.

We've provided the variables you should use for this model. The first CNN layer is a $k \times k$ convolution that takes in an image with one channel and outputs an image with c channels (where k is the value of `--cnn-n1-kernel` and c is the value of `--cnn-n1-channels`). The second conv layer is a $k_2 \times k_2$ convolution that takes in an image with c channels and outputs an image with 10 channels (where k_2 is the value of `--cnn-n2-kernel`). The output image has approximately half the height and half the width because of the stride of 2.

You can see that this model has 3 hyperparameters (k , k_2 , c) in addition to the learning rate, optimizer, batch size, and number of training iterations. Use the Jupyter notebook to help select a good number of channels for our first CNN layer.

1. Complete the `SimpleConvNN` model in `models.py`, by implementing the forward function. Remember, the output of this model should be logits! (See Figure 2) You should apply a ReLU activation to the output of the 2 convolutional layers. Implement this model to feed the output of the first convolution into the second convolution. Then, you should have an output which is of size `[batch_size, 14, 14, 10]`. To convert this into logits, you should use a max-pooling layer (for example, the one here <https://pytorch.org/docs/stable/nn.functional.html#max-pool2d>) over the channels. This will give us a vector of 10 items, each the max value for it's respective channel in the previous output.
2. Complete sections 2.2 and 3.2 of the notebook using this model. You should use these to help you select good parameters for this model. Namely, this model will

take a lot more steps to train than the previous model. Look at the graph of the dev loss over training steps. Does it seem like it could still be trending upward? That means you need to train longer! Try to find a setting that at least matches your ff model on dev data.

3. Save your best CNN model as `cnn.torch` and the test predictions it makes as `cnn-predictions.txt`. **Submit both the model file and the predictions file.**

1.8 Best Model (30 points)

In this section your goal is to maximize performance. Do whatever you'd like to your network configurations to maximize validation accuracy. Feel free to vary the optimizer, mini-batch sizes, the number of layers and/or the number of hidden units / number of filters per layer; include dropout if you'd like do; etc. You can even go the extra mile with techniques such as data augmentation, where input images may be randomly cropped and/or translate and/or blurred and/or rotated etc. We've added the `scikit-image` (<https://scikit-image.org/>) package to the `requirements.txt` file, if you want to take this approach.

However, there are a couple limitations: You may not add additional training data, and you must be able to store your model in a `.torch` file no bigger than 1MB.

Hints. You may want to focus on convolutional networks, since they are especially well suited for processing images. You may often find yourself in a situation where training is just *too slow* (for example where validation accuracy fails to climb after a 5 minute period). It is up to you to cut experiments off early: if training in a reasonable amount of time isn't viable, then you can try to change your network or hyperparameters to help speed things up. In addition, earlier we repeatedly asked that you vary other parameters as necessary to maximize performance. There is obviously an *enormous* number of possible configurations, involving optimizers, learning rates, mini-batch sizes, etc. Our advice is to find settings that consistently work well across simple architectures, and to only adjust these settings if insights based on training and validation curves suggest that you should do so. In particular, remember that Adam helps you avoid manual learning-rate tuning, and remember that very small minibatches and very large minibatches will both lead to slow performance, so striking a balance is important.

1. Complete the `BestNN` model in `models.py`. You need to define the features for this model, as well as the forward function.
2. Train this network. Remember, it's up to you to ensure that you can train your model in a reasonable amount of time!
3. Complete sections 2.3 and 3.3 of the notebook using this model. As always, you should use these cells to give you information about which values you should use for certain hyperparameters of your model, and to help you debug your model.
4. Save your best model as `best.torch` and it's test predictions as `best-predictions.txt`. **Submit both the model file and the predictions file.**

To receive full credit, you need to achieve a test accuracy which indicates that you have put a sufficient amount of effort into building a good classifier. This should improve over the best model in previous sections by a reasonable amount.

2 What to Submit

In this assignment you will submit two things.

1. **Submit your code (.py files), models (.torch files), and test predictions (predictions.txt files) to cs475.org as a zip file.**

Your .torch models *must be named* `ff.torch`, `cnn.torch`, `best.torch`. When the code loads a saved torch model, it is expecting the respective class to be unchanged, so **please make sure that the version of code you submit contains the correct version of the class for each model, so that each model file can be loaded without error**. You can test this by running `main.py` in test mode on each of the models.

The prediction files should similarly be named `ff-predictions.txt`, `cnn-predictions.txt`, `best-predictions.txt`.

Your code must be uploaded as code.zip with your code, models, and predictions in the root directory. By ‘in the root directory,’ we mean that the zip should contain `*.py` at the root (`/*.py`) and not in any sort of substructure (for example `hw1/*.py`). One simple way to achieve this is to zip using the command line, where you include files directly (e.g., `*.py`) rather than specifying a folder (e.g., `hw1`):

```
zip code.zip *.py *.torch *-predictions.txt
```

2. **Submit your Jupyter notebook to gradescope.com.**

Create a PDF file containing your notebooks (e.g. File → Download As → PDF via LaTeX (.pdf)). Be sure you have run the entire notebook so we can see all of your output. You will submit this to the assignment called “Homework 3: Deep Learning: Notebook”.

For the loss and accuracy graphs, please make sure that the graph you are plotting for each cell is the graph of the metrics for the model that **you are submitting**.

It is imperative that you submit all the requested files named as specified. After you upload code.zip, please download it to verify its contents.

You will need to create an account on gradescope.com and signup for this class. The course is <https://gradescope.com/courses/21552>. Use entry code M6ZX2X. See this video for instructions on how to upload a homework assignment: https://www.youtube.com/watch?v=KMPoby5g_nE.

3 Questions?

Remember to submit questions about the assignment to the appropriate group on Piazza: <https://piazza.com/class/jkqbzabvyr15up>.