


# From Regression to Classification with Logistic Regression

Mark Dredze

Machine Learning  
CS 601.475



## Classification

- Data  $\{(\mathbf{x}_i, y_i)\}_{i=1}^N$   $\mathbf{x}_i \in \mathcal{X}^M$   $y_i \in \mathcal{Y}$
- Learn: a mapping from  $\mathbf{x}$  to discrete value  $y$ 
  - $f(\mathbf{x}) = y$
- Examples
  - Spam classification
  - Document topic classification
  - Identifying faces in images

## Binary Classification

- We'll focus on binary classification
  - $y_i \in \{0, 1\}$
- Usually easy to generalize to multi-class classification

## Different Definition

### Fitting a function to data

- Fitting: Optimization, what parameters can we change?
- **Function: Model, loss function**
- Data: Data/model assumptions? How we use data?
- ML Algorithms: minimize a function on some data

## Evaluation

- Accuracy

$$\frac{\text{number of correct predictions}}{\text{total number of predictions}}$$


- Other measurements appropriate for some tasks
  - Ex. we care more about certain types of mistakes

## Regression


- Least squares regression
  - Outputs real number for each example
- It seems that classification should be easier!
- Let's use regression for classification
  - Learn least squares regression model  $f_w(x)=y$ 
    - $f_w(x)=w^T x$
    - If  $y>0$ , predict "True (1)"
    - If  $y\leq 0$  predict "False (0)"

## Regression for Classification

- $f_w(x)=0$  partitions the input space into two class specific regions
  - Linear decision boundary



Good



Bad

Figures by Tommi Jaakkola

## Regression for Classification

- Mismatch between regression loss and classification
  - Classification: accuracy
  - We don't care about large vs. small values of output
- Outliers problematic
  - Prediction of 42 for example is fine for classification, bad for regression
- We need output to be either 1 or 0

# Machine Learning

## Fitting a function to data


- Fitting: Solve for  $w$  given  $y$  and  $x$
- Function: Regression uses squared loss
  - Bad match for our task!
- Data: assume dependent variable linear combination of independent variables
- Our loss function doesn't match classification goals

## Logistic Function

- Quick fix: apply a function to the output of regression that gives desired value
- Logistic function


- Outputs between 0 and 1
- Scaling parameter  $\alpha$
- Most outputs are close to 1 or 0

$$g_{\alpha}(x) = \frac{1}{1 + e^{-\alpha x}}$$



## Logistic Function

$$g_{\alpha}(x) = \frac{1}{1 + e^{-\alpha x}}$$



## Logistic Regression

- We can combine the logistic function and our regression model

$$g(w^T \cdot x_i) = \frac{1}{1 + e^{-w^T \cdot x_i}}$$

- Notice: as  $w^T \cdot x_i$  becomes:
  - Large- output closer to 1
  - Small- output closer to 0

## Probabilistic View

- We want to model the probability of a label given the example
- Conditional likelihood  $p(y|x)$
- Consider
  - We could maximize the joint:  $p(y, x)$
  - Which can be factored as:  $p(x|y)p(y)$
  - However, the best label  $y$  is the same under both:
$$\arg \max_{y=0,1} p(x|y)p(y) = \arg \max_{y=0,1} p(y|x)$$
- Because  $x$  is fixed/given

## Probabilistic View

- We can now write the distribution as
$$p_w(y=1|x) = \frac{1}{1 + e^{-w^T \cdot x}}$$
- Which implies that
$$p_w(y=0|x) = \frac{e^{-w^T \cdot x}}{1 + e^{-w^T \cdot x}}$$
- The odds of the event is then
$$\frac{p_w(y=1|x)}{p_w(y=0|x)} = e^{w^T \cdot x}$$
- And the log-odds are
$$\log \frac{p_w(y=1|x)}{p_w(y=0|x)} = w^T \cdot x$$

## Generalized Linear Models

- Generalized linear models
  - A linear model whose output is passed through non-linear function
- Decision boundary/surface
  - An  $n-1$  dimensional hyper-plane that separates the data into two groups
  - The non-linearity gives a classification boundary
- We still have
  - Convex model
  - Hypothesis class: linear decision boundaries

## Logistic Regression Decisions

- Given parameters  $w$ , how do we make predictions?

$$p_w(y=1|x) = \frac{1}{1 + e^{-w^T \cdot x}}$$

- If output  $> .5$ , predict 1, else predict 0
- In addition to prediction, we have confidence in prediction
  - Confidence is the probability of the prediction

# Logistic Regression

## Fitting a function to data

- Fitting: Solve for  $w$  given  $y$  and  $x$
- Function: Generalized linear function: logistic over regression
- Data: assume dependent variable linear combination of independent variables

Squared loss (green): we told linear regression to predict 1 for positive class  
 Logistic loss (red): loss is suffered whenever  $p(y=1|x) < 1$

The loss suffered (y-axis) by predicting  $w^T \cdot x$  (x-axis) when positive is correct answer



## Objective Function: Likelihood

- Conditional data likelihood

$$p(Y|X, w) = \prod_{i=1}^n p(y_i | x_i, w)$$

## Conditional Log Likelihood

$$p(Y|X, w) = \prod_{i=1}^n p(y_i | x_i, w)$$

$$\ell(Y, X, w) = \log p(Y|X, w) = \sum_{i=1}^n \log p(y_i | x_i, w)$$

$$p(y=1|x, w) = \frac{1}{1 + e^{-w^T \cdot x}} \quad p(y=0|x, w) = \frac{e^{-w^T \cdot x}}{1 + e^{-w^T \cdot x}}$$

## Logistic Regression

### Fitting a function to data


- **Fitting: Solve for  $w$  given  $y$  and  $x$**
- Function: Generalized linear function: logistic over regression: conditional log likelihood
- Data: assume dependent variable linear combination of independent variables

## Function Optimization


- We have a function and want to maximize/minimize it
- How do we find the point at which the function reaches its max/min?

## Function Optimization

- Take the derivative, set it equal to 0, solve!
- Will this v



## Convex Functions



## Maximum Likelihood Estimation

- MLE: Find the value at which the likelihood is maximized
  - We'll talk about other options later in the semester
- Given the conditional log likelihood
  - Take the derivatives for parameters  $w$
  - Set each derivative to 0
  - $M$  equations and  $M$  variables
  - Solve for  $w$
- Problem
  - No closed form (analytical) solution for  $w$


## Convex Optimization

- The conditional maximum likelihood is concave
  - There is a single maximal solution
- We can maximize using convex optimization techniques
  - Its easy to optimize convex functions
  - There are **many** convex optimization algorithms

## Gradient Descent

- First order method: needs first order derivatives
- Assuming  $F(x)$  is defined and differentiable, then  $F(x)$  decreases fastest if we go from  $x$  in the direction of the gradient of  $F$ 
  - $-\nabla F(x)$ - vector of partial derivatives of  $F$
  - $x' = x - \gamma \nabla F(x)$  - Update
- For sufficiently small values of  $\gamma$ , the value of the function will get smaller

## Gradient Descent





## Derivatives

Objective:  
conditional log likelihood  $\ell(Y, X, w) = \log p(Y | X, w) = \sum_{i=1}^n \log p(y_i | x_i, w)$

Given the sigmoid as  $h_w(x) = \frac{1}{1 + e^{-w^T \cdot x}}$

We can rewrite compactly  $p(y | x) = (h_w(x))^y (1 - h_w(x))^{1-y}$

New objective  $\ell(Y, X, w) = \sum_{i=1}^N \log \{ (h_w(x_i))^y (1 - h_w(x_i))^{1-y} \}$

Derivative  $\frac{\partial \ell(Y, X, w)}{\partial w} = \sum_{i=1}^N (y_i - h_w(x_i)) x_i$

- The derivative is 0 when  $y_i = p(y_i | x_i, w)$
- Maximizing likelihood = minimize logistic error

## Gradient Descent Solution

$$w^{(t+1)} = w^t + \gamma \frac{\partial \ell(Y, X, w)}{\partial w}$$

$$w^{t+1} = w^t + \gamma \sum_{i=1}^N (y_i - h_w(x_i)) x_i$$

## Algorithm: Logistic Regression

- Train: given data X and Y
  - Initialize w to starting value
  - Repeat until convergence
    - Compute the value of the derivative for X, Y and w
    - Update w by taking a gradient step
- Predict: given an example x
  - Using the learned w, compute  $p(y | x, w)$

$$p(y=1 | x, w) = \frac{1}{1 + e^{-w^T \cdot x}}$$

- Note: many other optimization routines available

## Gradient Based Optimization


- Multiple methods available for optimizing the same objective function
  - First order methods
  - Second order methods
  - Adaptive methods
  - ...



## Alternate Methods

- Batch gradient descent
  - Utilize the gradient of all the data
  - Slow: need to consider all the data before making a single update

## Gradient Descent




## Stochastic Updates

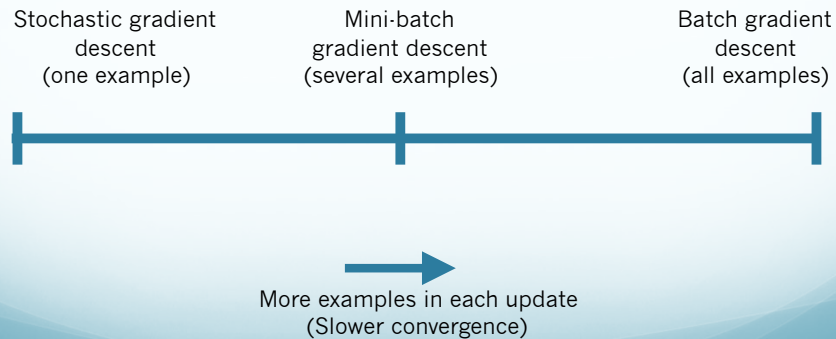
- Compute the gradient on a single example at a time

$$w^{t+1} = w^t + \gamma(y_i - h_w(x_i))x_i$$

## Gradient Descent



## Update Frequency



## Regularization

- Same over-fitting problems as least squares
- Add regularization term to objective to favor different considerations
- Similar options
  - Quadratic regularization (L2)
  - L1 regularization (sparse solutions)
- For each regularization optimize new objective function

## Summary

- Logistic regression
  - Learn  $p(y|x)$  directly with functional form of distribution
  - Maximize the data conditional log-likelihood
  - Equivalent to linear prediction
    - Decision rule is a hyper-plane
  - Regularization to prevent over-fitting