

# MRFs and Inference in Graphical Models

Mark Dredze



Machine Learning  
CS 601.475

## Outline

- Representation
  - What is a graphical model?
  - What does it represent
  - Conditional Independence
  - **Types of probabilistic models**
- Inference
  - **How can we compute probabilities?**
  - **Message Passing**
- Examples
  - Learning and inference


## Graph Types

- Edge type determines graph type
- Directed graphs
  - Edges have directions ( $A \rightarrow B$ )
  - Assume DAGs (no cycles)
  - Typically called Bayesian Networks
    - Popular in AI and stats
- Undirected graphs
  - Edges don't have directions ( $A - B$ )
  - Typically called Markov Random Fields (MRFs)
    - Popular in physics and vision



## Undirected Networks

- What happens when we use undirected edges?




- Called undirected graphical models
  - Markov Random Fields, Markov Networks

## Conditional Independence

- D-separation used directionality of the edges
- Can we define it without directionality?
  - Yes!
- If all paths from set A to set B pass through set C, then A is d-separated from B given C
  - Notice no distinction between head to head and tail to tail
  - “Explaining away” not an issue since no causation
  - Actually easier to check 😊

## Markov Blanket

- The absence of “explaining away” makes the Markov blanket simple as well
- The Markov blanket of a node contains the neighbors of the node



## Factorization

- How can we express the joint distribution as a product of functions over local sets of variables
  - The directions in directed graphs indicated conditional relationship
- Step 1 is easier than with directed graphs
  - Given two nodes  $x_i$  and  $x_j$ , they are conditionally independent given the entire graph if they are not neighbors


$$p(x_i, x_j | \mathbf{x}_{\setminus \{i,j\}}) = p(x_i | \mathbf{x}_{\setminus \{i,j\}}) p(x_j | \mathbf{x}_{\setminus \{i,j\}})$$

## Defining Factors

- The factorization of the joint distribution is such that  $x_i$  and  $x_j$  do not appear in the same factor
- Graph concept: clique
  - A set of nodes that are fully connected
    - There exists an edge between every pair of nodes
  - Maximal clique
    - A clique such that adding any other node means it is no longer a clique

## Cliques

- Cliques in the graph
  - A/B, A/C, B/D, B/C, C/D
  - Maximal cliques: A/B/C, B/C/D
  - A/B/C/D is not a clique since no edge from A to D
- We just need to use maximal cliques, since they contain all other cliques



## Factorization

- Define
  - $x_c$  as all nodes in clique C
  - $\psi_c(x_c)$  is a potential function over clique C
- We can define the joint distribution of the graph as a product of potential functions over maximal cliques

$$p(x) = \frac{1}{Z} \prod_c \psi_c(x_c)$$

- Cliques have taken the place of nodes and CPTs


## Partition Function

$$p(x) = \frac{1}{Z} \prod_c \psi_c(x_c)$$

- Notice that  $\psi_c(x_c)$  is not a probability
  - Must be  $\psi_c(x_c) \geq 0$
  - Will not sum to 1
- Therefore we need to normalize to get a probability
 
$$Z = \sum_x \prod_c \psi_c(x_c)$$
- Z is called the partition function

## Example


- A model where we have label Y and example X



- Test time, no Y
  - Estimate Y using X
- What model is this?

## Logistic Regression

- No generative story
  - Graph does not encode causality
  - We can say that X and Y are related
- Learning
  - We observe X and Y, maximum likelihood solution
- Prediction
  - Compute most likely value for Y given X




## Factorization

$$p(y | x) = \frac{1}{Z} \prod_m \psi_m(x_m, y)$$


$$\psi_m(x_m, y) = \exp(w_m \cdot f(x_m, y))$$

$$Z = \sum_y \prod_m \psi_m(x_m, y)$$



## Potential Functions

- The parameters correspond to potential functions




Y/X	0	1
0	.45	3
1	5	.2

M Tables

## Learning

- We assumed both examples (X) and labels (Y) for learning logistic regression
- Maximum likelihood solution
  - Values of potential functions learned using convex optimization

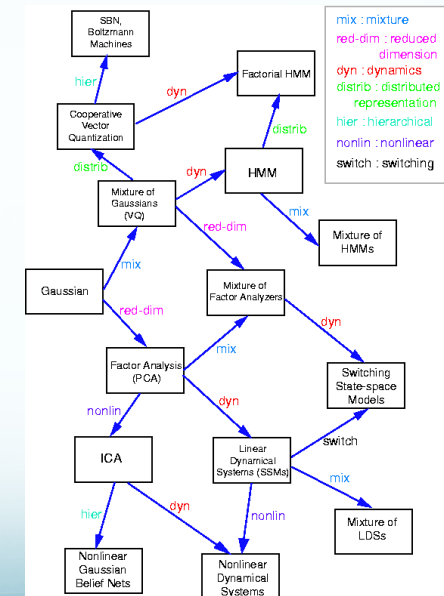


## MRF: Pros and Cons

- Pros
  - Define arbitrary potential functions
  - Much more flexibility than directed models
  - Easier to compute condition independence
  - Don't need to express causation relationship
- Cons
  - Z!!!
  - Sum over all states  $x$
  - M discrete nodes, each with K discrete states,  $K^M$ 
    - However, we just need Z for learning
    - To evaluate we need most likely option, so Z cancels

## Generative Model for Generative Models

There are tons of graphical models!



## Inference in Graphical Models


## Inference

- Computing probabilities of network configurations
  - We know some values of the network (observed)
  - How do we compute the posterior of a set of nodes?
- Previously, we did this by explicitly working out the probabilities
- How can we do this in an efficient and general way?

## Two Approaches

- Exact inference
  - We get the exact value of the probability we want
  - While some efficient algorithms exist, very slow for some graphs
- Approximate inference
  - Compute an approximation of the desired probability
  - The only solution for some types of graphs

## Chain Graphical Model



- Consider a linear chain of random variables
  - Notice this is undirected
  - We can convert directed models to undirected models
    - See book
- Joint distribution of the chain
 
$$p(\mathbf{x}) = \frac{1}{Z} \psi_{1,2}(x_1, x_2) \psi_{2,3}(x_2, x_3) \dots \psi_{N-1,N}(x_{N-1}, x_N)$$
  - Given N nodes with K states
  - Each potential function is a K\*K table
  - Joint has (N-1)K<sup>2</sup> parameters

## Chain Inference

- What is  $p(x_n)$  for some node n in the chain?
  - Assuming no observed nodes
- Sum over all other nodes in the chain
 
$$p(x_n) = \sum_{x_1} \dots \sum_{x_{n-1}} \sum_{x_{n+1}} \dots \sum_{x_N} p(\mathbf{x})$$
  - Notice there are K<sup>N</sup> values to consider in the summation
  - Our computations are **exponential** in the length of the chain

## Conditional Independence

- Conditional independence to the rescue!
  - We can write the joint in terms of potentials
  - Each potential depends on 2 nodes
- Plug the factorized joint into the marginal for  $p(x_n)$

$$p(\mathbf{x}) = \frac{1}{Z} \psi_{1,2}(x_1, x_2) \psi_{2,3}(x_2, x_3) \dots \psi_{N-1,N}(x_{N-1}, x_N)$$

$$p(x_n) = \sum_{x_1} \dots \sum_{x_{n-1}} \sum_{x_{n+1}} \dots \sum_{x_N} p(\mathbf{x})$$

## Conditional Independence

- Consider the final summation for variable  $x_N$
- Only one potential depends on  $x_N$
- We can perform this summation first to give a function of  $x_{N-1}$

$$\sum_{x_N} \psi_{N-1,N}(x_{N-1}, x_N)$$

- Since no other terms depend on  $x_N$  we can push this summation all the way in
- The same is true starting from the other side of the chain for  $x_1$

## Grouping Potentials

$$p(x_n) = \frac{1}{Z} \underbrace{\mu_\alpha(x_n)}_{\left[ \sum_{x_{n-1}} \psi_{n-1,n}(x_{n-1}, x_n) \cdots \left[ \sum_{x_2} \psi_{2,3}(x_2, x_3) \left[ \sum_{x_1} \psi_{1,2}(x_1, x_2) \right] \cdots \right] \right]} \underbrace{\mu_\beta(x_n)}_{\left[ \sum_{x_{n+1}} \psi_{n,n+1}(x_n, x_{n+1}) \cdots \left[ \sum_{x_N} \psi_{N-1,N}(x_{N-1}, x_N) \right] \cdots \right]}$$


- Cost?  $N-1$  summations over  $K$  states
  - Each computation is a factor with a  $K \times K$  table
  - $O(NK^2)$  – linear in the length of the chain

## Rewriting As Factors

- The marginal can be written in terms of two factors

$$p(x_n) = \frac{1}{Z} \mu_\alpha(x_n) \mu_\beta(x_n)$$

- Each factor depends only on the nodes to one side
- This information is passed along the network to  $x_n$
- We call this a message
  - Each message contains  $K$  values (for every  $x_n$ ),



## Calculating the Message

- Messages are computed recursively (based on other messages)

$$\begin{aligned} \mu_\alpha(x_n) &= \sum_{x_{n-1}} \psi_{n-1,n}(x_{n-1}, x_n) \left[ \sum_{x_{n-2}} \cdots \right] \\ &= \sum_{x_{n-1}} \psi_{n-1,n}(x_{n-1}, x_n) \mu_\alpha(x_{n-1}) \end{aligned}$$

- Base case

$$\mu_\alpha(x_2) = \sum_{x_1} \psi_{1,2}(x_1, x_2)$$

- The same is true for  $\mu_\beta(x_n)$



## Normalization Constant

- Z is a sum over all states in

$$p(\mathbf{x}_n) = \frac{1}{Z} \mu_\alpha(\mathbf{x}_n) \mu_\beta(\mathbf{x}_n)$$

- which is  $O(K)$
- This is an example of a message passing algorithm

## Computing All Marginals?

- What if we wanted to compute  $p(\mathbf{x}_n)$  for all  $n$ ?
  - Just run the message passing algorithm ( $O(NK^2)$ )  $n$  times:  $O(N^2K^2)$
- This is wasteful since we keep computing the same messages many times
  - Instead, compute them once and save them:  $O(NK^2)$

## Computing Probabilities

- What about observed variables?
  - Clamp their value, remove the sum
- What about joint distribution over two neighboring nodes?

$$p(\mathbf{x}_{n-1}, \mathbf{x}_n) = \frac{1}{Z} \mu_\alpha(\mathbf{x}_{n-1}) \psi_{n-1,n}(\mathbf{x}_{n-1}, \mathbf{x}_n) \mu_\beta(\mathbf{x}_n)$$

- This is great because we can parameterize the potentials and learn them
  - How do we learn these parameters?
  - EM! In fact, this is the E step (assume you have observations, compute the probability given the parameters)

## Factor Graphs


- Both types of graphical models (directed/undirected) allow a global function to be expressed as a product of local factors
  - This factorization is what makes them powerful
- We can make this explicit by representing these factors directly
- Write the joint distribution directly in terms of factors
  - $\mathbf{x}_s$  - a subset of the variables
  - $f_s$  - function of the subset of variables

$$p(\mathbf{x}) = \prod_s f_s(\mathbf{x}_s)$$



## Factor Graphs

- We can construct a factor graph to represent a given graphical model
  - For an undirected model, the factors are potential functions over the maximal cliques



## Factor Graphs

- Factor graphs are bipartite
  - Two types of nodes: nodes and factors
  - Nodes only link to factors, factors only link to nodes
- Multiple factor graphs represent the same distribution
  - Factor graphs can be more specific about factorization

## Sum Product Algorithm

- Factor graphs are used to derive the sum product algorithm
  - A powerful class of efficient exact inference algorithms for tree-structured graphs

- How do we find the marginal

$$\rho(x) = \sum_{\mathbf{x} \setminus x} \rho(\mathbf{x})$$


- $\mathbf{x} \setminus x$ : the set of variables in  $\mathbf{x}$  with variable  $x$  omitted
- We will replace  $p(\mathbf{x})$  with factor graph to obtain efficient algorithm

## Sub-Tree

- Let's consider a fragment of the tree, which allows us to partition the joint into two sets of factors

- We can write the joint as

$$\rho(\mathbf{x}) = \prod_{s \in \text{ne}(x)} F_s(x, X_s)$$



- $\text{ne}(x)$ : factors that are neighbors of  $x$
- $X_s$ : all variables in sub-tree connected via factor  $f_s$
- $F_s(x, X_s)$ : product of all factors associated with  $f_s$

## Factor Messages

- Combining these equations

$$\begin{aligned} \rho(\mathbf{x}) &= \prod_{s \in \text{ne}(x)} F_s(x, X_s) \\ \rho(x) &= \sum_{\mathbf{x} \setminus x} \rho(\mathbf{x}) \end{aligned} \quad \Rightarrow \quad \begin{aligned} p(x) &= \prod_{s \in \text{ne}(x)} \left[ \sum_{X_s} F_s(x, X_s) \right] \\ &= \prod_{s \in \text{ne}(x)} \mu_{f_s \rightarrow x}(x) \end{aligned}$$


- Messages between factor  $f_s$  to node  $x$

$$\mu_{f \rightarrow x}(x) = \sum_{X_s} F_s(x, X_s)$$

- The marginal is a product of all incoming messages

## Sub-Factors

- Notice that the factors themselves can be broken into factor sub-trees
- Each message is written



$$\begin{aligned} \mu_{f_s \rightarrow x_m}(x_m) &= \sum_{x_1} \dots \sum_{x_m} f_s(x, x_1 \dots x_m) \prod_{m \in \text{ne}(f_s) \setminus x} \sum_{X_{sm}} G_m(x_m, X_{sm}) \\ &= \sum_{x_1} \dots \sum_{x_m} f_s(x, x_1 \dots x_m) \prod_{m \in \text{ne}(f_s) \setminus x} \mu_{x_m \rightarrow f_s}(x_m) \end{aligned}$$

- $G_m$  is the product of all factors associated with  $x_m$
- Can replace last term by a message from node to factor

## Node Messages

- Messages between factor nodes and factors

$$\mu_{x_m \rightarrow f_s}(x_m) = \sum_{X_{sm}} G_m(x_m, X_{sm})$$

- Two different types of messages

- Nodes to factors
- Factors to nodes

- Messages passed by a node are a product of variables connected to that node


## Evaluating Messages

- How do we evaluate messages from variable nodes to factor nodes?

- Sub Graph factorization!

$$G_m(x_m, X_{sm}) = \prod_{l \in \text{ne}(x_m) \setminus f_s} F_l(x_m, X_{ml})$$

- $G_m(x_m, X_{sm})$  - associated with node  $x_m$
- $F_l(x_m, X_{ml})$  - associated with factor  $f_l$  linked to  $x_m$
- Product over all neighbors except  $f_s$



## Messages

- Using this definition we can write messages from nodes to factors

$$\begin{aligned}\mu_{x_m \rightarrow f_s}(x_m) &= \prod_{l \in ne(x_m) \setminus f_s} \left[ \sum_{x_{ml}} F_l(x_m, x_{ml}) \right] \\ &= \prod_{l \in ne(x_m) \setminus f_s} \mu_{f_l \rightarrow x_m}(x_m)\end{aligned}$$

- If only two neighbors, then just pass the message
- Requires input from all other factors before can send a message

## Recursion

- These messages are all based on recursion
- Base cases (leaves of the network)?

$$\mu_{x \rightarrow f}(x) = 1 \quad \mu_{f \rightarrow x}(x) = f(x)$$

## Summary: Sum Product

- We want to evaluate marginal  $p(x)$  in a tree factor graph
  - Assume  $x$  is root, start messages at the leaves
  - Propagate messages until  $x$  receives messages
    - Wait until all neighbors except one send message, then create new message
  - When  $x$  receives all messages, compute  $p(x)$


$$p(x) = \prod_{s \in ne(x)} \mu_{f_s \rightarrow x}(x)$$

- As before, we can compute the messages over the whole graph and use them to compute every marginal

## Normalization

- If this is an undirected graph, then the result is a distribution that is not normalized
  - If we started from a directed graph, it is normalized
- Normalization is easy
  - Recall with chains we computed  $Z$  by summing over all configurations for the node(s) of interest
  - We have the messages to compute this here

## Sum Product Example



## Messages


$$\mu_{f \rightarrow x}(x) = \sum_{X_s} F_s(x, X_s)$$

$$\mu_{f \rightarrow x}(x) = f(x)$$

$$\mu_{x_m \rightarrow f_s}(x_m) = \prod_{l \in ne(x_m) \setminus f_s} \mu_{f_l \rightarrow x_m}(x_m)$$

$$\mu_{x \rightarrow f}(x) = 1$$

## Sum Product Example



## Inference in Chains

- The marginal can be written in terms of two messages

$$p(x_n) = \frac{1}{Z} \mu_\alpha(x_n) \mu_\beta(x_n)$$

- This is sum product!
  - We just need to add factors.

## Most Likely Configuration

- We know how to find the probability of configuration
- How do we find the most likely configuration?
  - Run the Sum Product algorithm for each marginal
  - Select the most probable value for each node
- Problem: this gives a node specific max probability
- We want most likely values for all of the nodes together

## Inference

- The goal of inference is to find the probability of a certain configuration
- Examples
  - The probability of a configuration  $\mathbf{x}$ ?
  - The probability of  $x_2=1$ ?
  - **The most likely configuration?**

## Max Product Algorithm

- Computes the configuration of the graph that gives the highest global probability
- Replace sum in Sum Product with max
 
$$\mathbf{x}^{\max} = \arg \max_{\mathbf{x}} p(\mathbf{x})$$
- Key trick in Sum Product
  - $ab + ac = a(b + c)$
- Key trick in Max Product
  - $\arg \max(ab, ac) = a \arg \max(b, c)$

## Moving Out Terms

- Write out the max terms
  - $\mathbf{x}^{\max} = \arg \max_{\mathbf{x}} p(\mathbf{x}) = \arg \max_{x_1} \dots \arg \max_{x_N} p(\mathbf{x})$
- For chain nodes, we write
 
$$\begin{aligned} \max_{\mathbf{x}} p(\mathbf{x}) &= \frac{1}{Z} \max_{x_1} \dots \max_{x_N} [\psi_{1,2}(x_1, x_2) \dots \psi_{N-1,N}(x_{N-1}, x_N)] \\ &= \frac{1}{Z} \max_{x_1} \left[ \psi_{1,2}(x_1, x_2) \dots \left[ \max_{x_N} \psi_{N-1,N}(x_{N-1}, x_N) \right] \dots \right] \end{aligned}$$

## Message Passing

- To avoid underflow, take  $\log(\max)$
- Using the message passing formulation from Sum Product, we obtain new messages for Max Product

$$\mu_{f_s \rightarrow x}(x) = \max_{x_1 \dots x_N} \left[ \log f_s(x, x_1 \dots x_N) + \sum_{m \in \text{ne}(f_s) \setminus x} \mu_{x_m \rightarrow f}(x_m) \right]$$

$$\mu_{x \rightarrow f_s}(x) = \sum_{l \in \text{ne}(x) \setminus f_s} \mu_{f_l \rightarrow x_m}(x)$$

- Base case

$$\mu_{f_s \rightarrow x}(x) = \log f_s(x)$$

$$\mu_{x \rightarrow f_s}(x) = 0$$


Root node

$$x^{\max} = \arg \max_x \left[ \sum_{s \in \text{ne}(x)} \mu_{f_s \rightarrow x}(x) \right]$$

## Messages from Root

- We can't pass messages back as before
  - There may be multiple configurations that have max value for  $p(\mathbf{x})$
  - Need to distinguish between these configurations
- Solution: keep track of which states correspond to the same max configuration
  - Store quantities:
 
$$\phi(x_n) = \arg \max_{x_{n-1}} [\log f_{n-1,n}(x_{n-1}, x_n) + \mu_{x_{n-1} \rightarrow f_{n-1,n}}(x_{n-1})]$$
    - The argument at  $x_{n-1}$  that gave the highest probability
    - To find the best path we follow the values for  $\phi(x)$

## Example for a Chain



## Inference Review

- Message Passing Algorithms for Exact Inference
  - Each node computes a local message to send to its neighbors
  - Sum Product
    - Compute marginals for nodes in the graph
    - Efficient way to compute all marginals
  - Max Product
    - Find the highest probability configuration
    - Back track through graph to decode entire configuration
- Non-tree structured graphs
  - Exact inference algorithms exist for certain types of graphs

# Graphical Models Review

## Graphical Models Overview

- Types of models
  - Directed Graphical Models
    - Bayesian Networks
  - Undirected Graphical Models
    - Markov Random Fields
  - Factor Graphs

## Model Types

- Directed Graphical Models
  - CPTs (everything is already normalized)
  - Training often by counting
- Undirected Graphical Models
  - Factors replace CPTs
  - Learn the factors by discriminative training
  - Generative training requires more complex  $Z$

## Efficient Algorithms

- Message passing algorithms
- Computing marginal probabilities
  - Sum Product Algorithm for trees
- Computing max probability configuration
  - Max Product Algorithm for trees
  -



## Learning Settings

- Unsupervised learning
  - We don't observe all of the variables
- Semi-supervised learning
  - We observe variables for only some examples
- Supervised learning
  - We observe all the variables at training

## Remaining Issues

- Approximate Inference
  - Computing probabilities in non-tractable distributions
    - Distribution is dimensionality of latent space is too high
    - Posterior distribution is complex and can't be computed analytically
- Examples
  - Variational inference
  - Sampling methods

## Variational Inference

- We can't compute the marginal  $p(\mathbf{X}, \mathbf{Z})$ 
  - Approximate  $p(\mathbf{Z}|\mathbf{X})$  and  $p(\mathbf{X})$
  - Maximize  $p(\mathbf{X})$ 
    - Decompose  $p(\mathbf{X})$  using EM idea into lower bound and KL
- Want to minimize KL between  $p$  and our estimate  $q$ 
  - $q$  is non-tractable
  - Replace with an approximating distribution  $q^*$

## Sampling Methods

- We want to compute the marginal for  $p(\mathbf{X})$ 
  - Cannot compute the expectation of the distribution
  - We can draw samples from the distribution
- Idea: approximate the expectation by taking many samples
- Examples
  - Rejection sampling
    - Sample and throw away points that don't match what you want
  - Gibbs Sampling (Markov Chain Monte Carlo)
    - Draw a new value for one variable, update the others

## Remaining Issues

- General graph structures
  - Not trees
- Generative training for undirected models
  - Tricks for computing and estimating  $Z$
- Learning graph structures
  - Structural EM
- Continuous distributions