


Structured Prediction




Binary Prediction

- So far we've focused on binary prediction
- Machine learning can do much more than that!
- Multi-class prediction
- Complex structured outputs

Sequence Prediction




<http://cs.wellesley.edu/~cs313/projects/project8/images/OneStrand.jpg>



Trees


In what city did Piotr's last 1st place finish occur?



$R[\text{Venue}].\text{argmax}(\text{Position.1st}, \text{Index})$


<http://nlp.stanford.edu/software/sempre/wikitable/images/piotr.png/>

Object Recognition



<https://3c1703fe8d.site.internapcdn.net/newman/gfx/news/hires/2014/googleteamri.png>

Ranking



What is Structured Prediction?

- Input: x
 - Typically a structured input
 - Maintain structure of input in x
 - Do not flatten into list of features in an instance
- Output: y
 - y is now from a large set of possible outputs
 - Outputs defined based on input
 - Often exponential in size of input

Previous Approaches

- Naturally multi-class algorithm
 - Neural networks
 - Decision Trees
- Reduction to binary
 - e.g. one classifier per class
- These methods don't work when
 - Exponential number of output
 - Outputs defined based on input

Structured Prediction Challenges

- Scoring
 - How do we assign a score/probability to a possible output structure
- Search/Inference
 - Find the best scoring output structure
 - How do we search through an exponential number of options?

Outline

- Graphical models for structured prediction
 - Sequences: HMMs and CRFs
- Score based linear models
 - Perceptron, SVM
- Deep networks

Graphical Models: Sequence Models

Sequential Events

- Many events happen in sequence
 - Weather on consecutive days
 - Words in a written sentence
 - Spoken sounds by a person
 - Movements in the stock market
 - DNA base pairs
- We want to model these sequences with a graphical model

Sequential Models


- Simple approach
 - Each event is independent



- $p(x_1, x_2 \dots x_N) = \prod_{n=1}^N p(x_n)$
- Simple, but not very helpful

Sequential Models

- Complex approach
 - Each event is dependent on previous events




- $p(x_1, x_2 \dots x_N) = \prod_{n=1}^N p(x_n | x_1 \dots x_{n-1})$
- Captures dependencies, but way too complex

Markov Assumption


- The current state depends on a fixed number of previous states
 - The weather today depends on the past three days, but NOT two weeks ago
 - The next word in the sentence depends on the past three words, but nothing before
- Pro: makes for simple models
- Con: doesn't capture full history

Markov Chains

- First order Markov chain




- Second order Markov chain




DNA Example

- Given a sequence of base pairs, find regions that encode for Genes




- What model is this?
- Hidden Markov model

Markov Blankets



- The Markov blanket for z_n contains z_{n-1} , z_{n+1} and x_n
- The Markov blanket for x_n contains z_n
- Nodes are dependent on a small number of neighbors

Conditional Probability Tables



- Tables are the same for each node

Joint Probability of HMM

- The joint probability of an HMM

$$p(\mathbf{X}, \mathbf{Z} | \theta) = p(z_1 | \pi) \left[\prod_{n=2}^N p(z_n | z_{n-1}, \mathbf{A}) \right] \left[\prod_{m=1}^N p(x_m | z_m, \phi) \right]$$

- \mathbf{A} - transition probabilities (matrix)
 - A_{ij} is the probability of moving from state i to j
- π - vector with starting probabilities
- ϕ – emission probabilities (matrix)
 - ϕ_{ij} is the probability of state i and emitting observation j

Unsupervised Training

- How do we train a probabilistic model?

- Maximum likelihood!

$$\max_{\theta} p(\mathbf{X}, \mathbf{Z} | \theta)$$

- Problem: we don't know \mathbf{Z}

- Solution: EM

- Step 1: Write the complete data likelihood

$$p(\mathbf{X} | \theta) = \sum_{\mathbf{Z}} p(\mathbf{X}, \mathbf{Z} | \theta)$$

$$p(\mathbf{X}, \mathbf{Z} | \theta) = p(z_1 | \pi) \left[\prod_{n=2}^N p(z_n | z_{n-1}, \mathbf{A}) \right] \left[\prod_{m=1}^N p(x_m | z_m, \phi) \right]$$

EM for HMMs

- E-Step
 - Find the expected values for the hidden variables \mathbf{Z} given the model parameters
 - The most likely \mathbf{Z} given \mathbf{X} and current model parameters
- M-Step
 - Pretend to observe the values for \mathbf{Z}
 - Update model parameters \mathbf{A} , π , ϕ to maximize complete data likelihood

EM for HMMs

- E-Step
 - Given Q function, evaluate probabilities for \mathbf{Z}

$$Q(\theta, \theta^{old}) = \sum_{\mathbf{Z}} p(\mathbf{Z} | \mathbf{X}, \theta^{old}) \log p(\mathbf{X}, \mathbf{Z} | \theta)$$

- For the HMM

$$Q(\theta, \theta^{old}) = \sum_{k=1}^K \gamma(z_k) \log \pi_k + \sum_{n=2}^N \sum_{j=1}^K \sum_{k=1}^K \xi(z_{n-1}, j, z_{nk}) \log A_{jk} + \sum_{n=1}^N \sum_{k=1}^K \gamma(z_{nk}) \log \phi_k$$

$$\gamma(z_n) = p(z_n | \mathbf{X}, \theta^{old}) \quad \pi, \phi \text{ and } \mathbf{A} \text{ are model parameters}$$

$$\xi(z_{n-1}, z_n) = p(z_{n-1}, z_n | \mathbf{X}, \theta^{old})$$

EM for HMMs

- How can we get these values?

$$\gamma(z_n) = p(z_n | X, \theta^{old}) \quad \xi(z_{n-1}, z_n) = p(z_{n-1}, z_n | X, \theta^{old})$$

- What is the probability of being in state z_n ?
- What is the probability of being in state z_n and z_{n+1} ?
- Inference
 - Sum Product Algorithm
 - Forward Backward in HMMs


EM for HMMs

- M Step
 - Maximize Q function with respect to π, Φ and A
- Assuming values for Z , we get

$$\pi_k = \frac{\gamma(z_{1k})}{\sum_{j=1}^K \gamma(z_{1j})} \quad A_{jk} = \frac{\sum_{n=2}^N \xi(z_{n-1,j}, z_{nk})}{\sum_{l=1}^K \sum_{n=2}^N \xi(z_{n-1,l}, z_{nl})} \quad \phi_{ik} = \frac{\sum_{n=1}^N \gamma(z_{nk}) x_{ni}}{\sum_{n=1}^N \gamma(z_{nk})}$$

Prediction

- Given a new sequence X , find the most likely set of states to have generated X
- Find the sequence Z with the maximum probability given X



Prediction

- How do we find the most likely state?

$$\arg \max_z p(z | X, \theta)$$
- Inference
 - Max Product Algorithm
 - Viterbi Decoding

Supervised Training


- We actually observe \mathbf{Z}
 - Just compute M step a single time
 - Very fast and easy
- What if we observe only some \mathbf{Z}
 - Case 1: only some examples are labeled with \mathbf{Z}
 - Case 2: each example has only some labels for \mathbf{Z}
- Semi-supervised Learning
 - Use EM algorithm but fix \mathbf{Z} when known

Notes on HMMs

- An HMM can have continuous or discrete emissions
 - Discrete- base pair, word in sentence
 - Continuous- stock price, frequency of a sound
- An HMM has discrete hidden states
- A Linear Dynamical System has continuous hidden states


Sequence Models

- An HMM is a directed graphical model
 - Bayesian Network



- What happens if we have an undirected graphical model?
 - Markov Random Field
 - Conditional Random Fields

Conditional Random Fields



- Replace conditional probabilities with potential functions
- The joint probability is a product of potential functions

$$p(\mathbf{X}, \mathbf{Z}) = \frac{1}{Z} \prod_A \psi_A(\mathbf{x}_A)$$

Joint Probability

- The joint probability of the HMM can be written using potential functions

- HMM

$$p(\mathbf{X}, \mathbf{Z} | \theta) = p(z_1 | \pi) \left[\prod_{n=2}^N p(z_n | z_{n-1}, \mathbf{A}) \right] \left[\prod_{m=1}^N p(x_m | z_m, \phi) \right]$$

- CRF

$$p(\mathbf{X}, \mathbf{Z} | \theta) = \frac{1}{Z} \psi(z_1) \left[\prod_{n=2}^N \psi(z_n | z_{n-1}) \right] \left[\prod_{m=1}^N \psi(x_m | z_m) \right]$$

$$p(\mathbf{X}, \mathbf{Z} | \theta) = \frac{1}{Z} \prod_A \psi_A(\mathbf{x}_A, \mathbf{z}_A) \quad \mathbf{Z} = \sum_{\mathbf{X}, \mathbf{Z}} \prod_A \psi_A(\mathbf{x}_A, \mathbf{z}_A)$$

Learning a CRF

- Given some data, we want to learn a CRF
- How should we learn the model?
 - Maximum likelihood!
 - Maximize the probability of the data given the model
- Questions
 - What are the parameters of our model?
 - The potential functions
 - What is the objective?
 - How do we compute model probabilities efficiently?

Model Parameters

- HMM: model parameters are CPTs
 - CRF: CPTs replaced with potential functions
- Parameterize the potential functions
 - Learn the parameters

$$\psi(\mathbf{x}, \mathbf{z}) = \exp \left\{ \sum_k \theta_k f_k(\mathbf{x}, \mathbf{z}) \right\}$$
 - Parameters θ determine value of potential function
 - f_k is a feature function
 - $f_k = 1$ if x is the base pair "G"
 - Notice: linear combination of features (linear model!)

Objective

- Let's maximize the likelihood of the data
 - For a single example

$$p(\mathbf{X}, \mathbf{Z} | \theta) = \frac{1}{Z} \prod_A \psi_A(\mathbf{x}_A, \mathbf{z}_A)$$
 - What about Z ?

$$Z = \sum_{\mathbf{X}, \mathbf{Z}} \prod_A \psi_A(\mathbf{x}_A, \mathbf{z}_A)$$
 - Sum over all X !
 - All possible sequences of base pairs
 - It's too hard to learn X

Discriminative Training

- Solution: don't learn $p(X)$!
- Maximize the conditional likelihood of the data
 - For a single example

$$p(\mathbf{Z} | \mathbf{X}, \theta) = \frac{1}{Z} \prod_A \psi_A(\mathbf{x}_A, \mathbf{z}_A)$$

$$Z = \sum_{\mathbf{z}} \prod_A \psi_A(\mathbf{x}_A, \mathbf{z}_A)$$

- CRF Conditional log likelihood of all examples

$$\log p(\mathbf{z} | \mathbf{x}) = \sum_{i=1}^N \sum_{t=1}^T \sum_{k=1}^K \theta_k f_k(z_{it}, z_{it-1}, \mathbf{x}_{it}) - \sum_{i=1}^N \log Z(\mathbf{x}_i)$$

Problems with Objective

- Recall for logistic regression (discriminative training) maximum likelihood over-fit the data
- Solution: regularization

$$\log p(\mathbf{z} | \mathbf{x}) = \sum_{i=1}^N \sum_{t=1}^T \sum_{k=1}^K \theta_k f_k(z_{it}, z_{it-1}, \mathbf{x}_{it}) - \sum_{i=1}^N \log Z(\mathbf{x}_i) - \sum_{k=1}^K \frac{\theta_k^2}{2\sigma^2}$$

- Gaussian prior ($\mu=0, \Sigma=\sigma^2 I$)

Training a CRF

- The conditional log likelihood is convex

- Take the derivative and solve for θ

$$\frac{\partial L}{\partial \theta_k} = \sum_{i=1}^N \sum_{t=1}^T f_k(z_{it}, z_{it-1}, \mathbf{x}_{it}) - \sum_{i=1}^N \sum_{t=1}^T \sum_{\mathbf{z}} f_k(\mathbf{z}, \mathbf{x}_i) p(\mathbf{z} | \mathbf{x}_i) - \sum_{k=1}^K \frac{\theta_k}{\sigma^2}$$

- The derivative is 0 when
 - The last term (regularizer) is 0
 - The first term and the second term cancel each other
 - First term: the expected value for f_k under the empirical distribution (from the data)
 - Second term: expectation for f_k given model distribution

Computing Probabilities

- What do we need to compute the values in the derivative?
 - Marginal probability of $p(\mathbf{z} | \mathbf{x}_i)$
 - The normalization constant Z
 - Total score for all possible labelings of the sequence
 - Sum Product Algorithm (Forward-Backward)
- Prediction
 - Sequence of states with max probability
 - Prediction
 - How do we find the highest probability sequence?
 - Max Product Algorithm (Viterbi decoding)

CRF Summary

- CRFs are
 - Markov Random Fields
 - The MRF equivalent of a supervised HMM
 - Discriminatively trained using conditional log likelihood
 - Linear models (recall linear potential functions)
- CRF training contains versions of
 - Sum Product Algorithm
 - Max Product Algorithm
 - Convex optimization

Why CRFs?

- CRF training is much harder than HMM
 - Computing gradients, optimization vs. counting
 - 11 labels, 200k tokens: 2 hours / 45 labels, 1m tokens: 1 week
- Why bother?
 - HMMs require
 - Assumptions of causation / generative story
 - Independence assumptions for observations
 - These aren't problems for CRFs!
 - Can allow arbitrary dependencies
 - Transition can depend on x and z
 - Can condition on the whole sequence x
 - Recall:
 - Generative models limit the features
 - Discriminative models can have any types of features

HMMs and CRFs

- Generative/Discriminative pairs
 - A generative and discriminative parametric model family that can represent the same set of conditional probability distributions
 - Naïve Bayes/Logistic Regression
 - HMM/CRF
- HMM is a Naïve Bayes classifier at each node
- CRF is a Logistic Regression classifier at each node

Score Based Methods


- We have several methods that produce a score for a label
 - Perceptron, SVM
 - Generalized linear models
 - Not probabilistic graphical models
- Can we use these for structured prediction?

Multi-class Encodings

- Encode multiple labels into a single weight vector
- Update the entire vector at once in a single model

1 of K Encoding

- Encode all K labels in a single linear classifier
 - Lengthen the weight vector to $M \times K$
 - A segment of length M learns weights for class k



1 of K Encoding for Perceptron

- Learning:
 - Given example x , get score for each possible label
 - If highest scoring label \neq correct label
 - Update
 - Highest scoring incorrect label
$$w_i' = w_i - f_i(\hat{y}, x)$$
 - Correct label
$$w_i' = w_i + f_i(y, x)$$
 - This takes a gradient step
 - Increases the score of the correct label *and* decreases the score of the incorrect label

1 of K Encoding for Perceptron

- Prediction:
 - Given example x
 - For each label y :
 - Generate each feature $f_i(y, x)$
 - This returns a feature vector, where only the parameters corresponding to y are non-zero
 - Assign score = $w \bullet x$
 - Return the label with the highest score

Beyond Multiclass

- We used a 1 of K encoding to extend Perceptron to multiclass
- What about for general structured problems?
 - Sequence
 - Ranking
 - Trees
- Can we learn this using a Perceptron style algorithm?
 - Yes!

Generalized Perceptron

- On each round
 - Receive example x
 - Predict $\hat{y} = \text{sign}(w \cdot x) \longrightarrow \hat{y} = \arg\max_{y \in L} w \cdot f(x, y)$
 - Receive correct label $y \in \{+1, -1\} \longrightarrow y \in L$
 - Suffer loss $\ell_{0/1}(y, \hat{y}) \longrightarrow \ell(y, \hat{y})$
 - Update $w \quad w^{i+1} = w^i + y_i x_i \longrightarrow w^{i+1} = w^i + f(y_i, x_i) - f(\hat{y}, x_i)$

Feature Functions

- How do we get this encoding?
- Feature functions!
 - Define a function $f_i(y, x)$
 - Returns the value of the i th feature based on the label and instance
 - Ex. Does this document contain the word “sports” AND is $y = \text{“travel”}$
 - Create a duplicate of each feature for each label to obtain a unique position in w

Structured Perceptron

- On each round
 - Receive example x
 - Predict $\hat{y} = \arg\max_{y \in L} w \cdot f(x, y)$ Need efficient procedure since L may be exponential
 - Receive correct label $y \in L$
 - Suffer loss $\ell(y, \hat{y})$ Loss sensitive to comparing structured objects
 - Update $w \quad w^{i+1} = w^i + f(y_i, x_i) - f(\hat{y}, x_i)$ General representations for each prediction

Example: Sequence Labeling

- Label a sequence of words with part of speech tags

John saw Mary with the telescope
noun verb noun preposition article noun

- We want to assign all tags for sentence at once

Requirements

- Procedure for finding best sequence using w
 - Yes! Viterbi decoding
 - We'll cover this in graphical models
- Loss sensitive to sequences
 - Yes! Hamming distance
 - How many labels disagree
- General representation
 - Features are based on sentence and label sequence

General Perceptron

- Perceptron updates can be used in any decision making process
- Requirements
 - Find best decision
 - Knowing incorrect decision
 - Featurize decisions

Other Types of Output

- Hierarchical Classification
 - Labels are organized in hierarchy
 - Mistaking "football" for "baseball" better than "football" for "NYSE"
- Sequences
 - Each label depends on neighboring labels
 - Part of speech
- Graphs
 - Output a graph or a tree given some input
 - Syntactic parse trees

One More Trick

- You can represent a structured Perceptron using a binary Perceptron

- Notice that:

- Structured form $w^{j+1} = w^j + f(y_i, x_i) - f(\hat{y}, x_i)$

$$x' = f(y_i, x_i) - f(\hat{y}, x_i)$$

$$y' = 1$$

- Binary form $w^{j+1} = w^j + y' x'$

- Prediction: Find highest scoring label according to w

- It will be greater than all other labels